

Arranging an Audio Track to other Genres by using CycleGAN-based Deep Learning Model *

Alex DongHyeon Seo
dseo22@wisc.edu

Hyecheol (Jerry) Jang
hyecheol.jang@wisc.edu

Stella Kim
ykim736@wisc.edu

Abstract

Changing the genre of a song is one of the methods used when composing music. To the best of our knowledge, musicians usually add their new ideas to the song, while trying to keep most of the special characteristics of the original song when arranging music. Similar to other artistic tasks that require human creativity, converting the genre of a song takes a significant amount of time and effort. In this project, we propose a method to translate a music genre by using machine-learning, which can generate a new song with comparably less amount of time than humans. Specifically, we utilized cycleGAN based model to translate a soundtrack to another soundtrack.

Due to the complexity and difficulties of dealing with audio data, our model is able to handle files written with MIDI (Musical Instrument Digital Interface) specification only with specific characteristics. In the near future, we expect to expand our project to use regular audio files rather than MIDI to do our tasks to generalize our model. By doing so, we hope our model to be used for the general public without further modifications.

1. Introduction

Recent advancements of Artificial Intelligence grant computers the abilities to mimic the noble creativity of the most intellectual lives. One of the advancements is computer-vision which allows computers to understand images just like human beings. However, unlike images, music has not been studied as extensive in the field of deep learning. Thus, we thought it would be interesting to do a project using data that represent audio.

Today, we are given many resources to share and access to all different types of music through apps like Spotify or YouTube. Due to those music apps and newly invented comfortable ear buds such as AirPods, many people seem

to be always listening to music on a daily basis. People listen to music while studying, driving, working and doing all other different types of tasks. Because of such increase in number of music listeners and people who are interested in music, musicians are always trying hard to provide new musical experience to the listeners. One of the ways that musicians create a song is changing a song to a different style of music while keeping the main characteristics of the original song. This is different from creating a totally new song from scratch and this way of changing a song to a new song by adding some new ideas is called music arrangement. Music arrangement can be done differently by making some different changes and one of the changes that could be made to a song is the genre. Only changing the genre of a song while keeping the main theme of the song might sound trivial but it is actually a very difficult task which takes quite a long time and much effort. Therefore, we would like to see if changing the genre of music could be done using machine-learning methods.

Our goal is very similar to AmazonWeb Service's DeepComposer¹, a keyboard that allows people to create a new song or arrange to a different genre with a simple combination of notes of their own. DeepComposer is based on deep learning models using generator and discriminator to update the music. Slightly different from DeepComposer, our project will focus on transferring the genre of the song to the user's desired genre.

2. Literature Review / Backgrounds

2.1. Audio Files

Our initial goal was to use raw audio files to train and test our model to generate a new song with a change in the genre of the original song.

However, while trying to find any good examples of machine learning models that use binary audio files for training and testing, we realized that there is not a well-made model that could be used for raw audio files. Unlike images, there

*Project proposal for Spring 2020, University of Wisconsin-Madison STAT453 Introduction to Deep Learning and Generative Models course (Instructor: Sebastian Raschka);
All authors contributed equally

¹AWS DeepComposer Product Description Page: <https://aws.amazon.com/deepcomposer/>

are not many models that are made for audio files. Thus, we decided to make a small change in our plan.

Instead of using audio files directly to the model, since we could not find a model that takes audio files as input, we thought we could use a converter that translates audio files to MIDI files. Even though we failed to find a machine learning model that uses audio files, we were able to find some models that use MIDI files. Because there are some models that were used with MIDI files, we thought it would work out if we can translate audio files to MIDI files and use them to train the model.

However, we faced another difficulties. We could not find the right converter that could be used in converting the type of file from audio such as MP3 to MIDI. After failing to find a way to audio file directly to a machine learning model and also failing to find a good converter that changes audio files into MIDI files, we decided to use MIDI files and build a model that takes in MIDI files as input for training and testing.

MIDI Files MIDI file is a type of file that represents audio using binary based on 128 representations of different instruments [1]. Inside a MIDI file, there are instructions written in binary that describe the attributes and characteristics of the song. For example, the file contains the information about the velocity of the song, which instruments to be played at which point of the track time and any other details that are needed to make the music.

Some advantages of using MIDI files is that MIDI files specify the 128 different instruments independently. As mentioned above, the MIDI files has their own way of representation of 128 instruments and they can be used in various ways depending how they are written in the MIDI file. Because those instruments are specifically described how they are played at which speed and time in MIDI files, every details of the information about the instruments that make up the sound of songs could be recorded. This is good for the purpose of our project since when generating a new song based on a song that already exists, it is very important to keep most of the details and the main style of the song.

Therefore, having the specifications of the music is helpful in maintaining the important sounds and notes when making a new song based off of the original version.

2.2. Model Architectures

2.3. Baseline Code

3. Proposed Method

4. Experiments

We generated two version of model to proof our idea. The first model utilizes the original soundtrack without any filtration or merging. The second model we tried merges

all the channels in MIDI file to one track (representing the piano's sound).

All the model follows the basic structures that proposed on Brunner et. al. [2]. For the second model, we tried four different settings during the training time. The first setting uses learning rate of 0.0002, epoch of 100 and second setting uses learning rate of 0.0001 and epoch of 200. The first and second setting uses 64 convolution filters on generator. For the third and fourth setting the learning rate and the epoch numbers are identical to the first and second setting, but it uses 128 convolution filters on generator.

4.1. Dataset

Finding a large amount of music having paired labels(the arranged/mixed music based on the same track) is very challenging and even considered as impossible task, since there is a limited number of arranged tracks for each music(with high possibility, there is no arranged track). To use cycle-GAN [3], we at least need to have unpaired labeled data. More specifically, we need music sources which have genre information as their label. Also, because the music industry is one of the most sensitive markets toward copyright issue, it is also important to only use the music tracks which do not have any restrictions on the usage for research purposes.

Data Collection For this project, we collected MIDI files of different genres, including Classical, Pop, Rock, and Jazz. We first tried collecting MIDI files of all four different genres from the same database to ensure the fairness and equality across all the files; however, we could not find a good database that has good MIDI files of all the four genres we intended to use for the project. Thus, we allowed ourselves to collect the files of each genre from different databases. At the end, we collected MIDI files of classical music from mfiles² while MIDI files of Pop and Rock are from midiworld³. However, we had some difficulties finding a good set of MIDI files of Jazz music so we used the same data with Brunner et al. [2]

For classical, pop and rock MIDI files, we did data crawling since those files are from websites. To do data crawling, we first used import.io,⁴ which is a web application that gets all the text contents of a webpage including links. Then we used R to cut off all the unnecessary text data we got from import.io to only select the download links for files. After cleaning up the data, then we used for-loop in R to go through the iterations and download all the files that are

²Classical Database: <https://www.mfiles.co.uk/classical-midi.htm>

³Pop Database: <https://www.midiworld.com/search/?q=classic>

Rock Database: <https://www.midiworld.com/search/?q=rock>

⁴Import.io product website: <https://www.import.io/>

linked to addresses that we found earlier. For Jazz, we simply downloaded the files from GitHub.

Dataset Analysis After we downloaded the MIDI files, we maneuver over the collected data to be more familiarized with those. At a glance, as the MIDI files are usually generated by the amateurs, rather than officially released from professional entertainment companies, they are not well-organized. For instance, Some of them contain meaningless trailing and ending whitespaces rather than their first note appears at time 0.

Moreover, as there exists the genres having characteristics from more than one genre (for instance, pop rock has characteristics from both pop and rock musics), it is hard to tell which musics belongs to which genre for some soundtrack. Also, the existence of variations in one genre (for example, inside of the rock genre, we have hard rock, progressive rock, punk rock and so on) made us hard to pinpoint the specific characteristics that each genre has. However, as there is no available databases such that categorizes the MIDI files in very specific genres, we decide to stay on using the big genre classes.

Data Pre-Processing Due to the unstructuredness of the retrieved dataset, we are forced to preprocess the data. During the data preprocessing, the baseline code dropped the songs which have time signature change, of which the first beat does not start at time zero, and of which time signature is not 4/4 [2]. The filter that Brunner et. al used is implemented in a very precise and specific way, which makes the songs that go through the filter ruins the song in such way that the music does not represent the intended genre anymore.

For the first experiment to make the model more generalized, we removed those filters before we processed the data. For the second experiment, which merges all the sound channels into one channel, we removed the trailing and ending whitespaces so that it can minimize the blank audio slices after cutting the sound to fixed-length slices.

4.2. Software

The baseline code was written in TensorFlow 1.4, which is very out-dated. We updated code so that it is at least runs without any warning or errors on the latest TensorFlow 1.x. As of now, though the TensorFlow 1 is outdated, we decided to use TensorFlow 1 as we are more familiar to.

To read and manipulate MIDI files properly, we utilizes `pretty_midi`⁵ and `Pypianoroll`⁶.

⁵`pretty_midi` Documentation: <http://craffel.github.io/pretty-midi/>

⁶`Pypianoroll` Documentation: <https://salu133445.github.io/pypianoroll/>

4.3. Hardware

For local testing before uploading the code and start train our model, we used Dell XPS 9550, which equipped Intel's i7-6700HQ @ 2.6Ghz, DDR4 2133Mhz, and NVIDIA GeForce GTX 960M with vRAM of 2GB.

After testing locally, for our main training, we used Google's Colaboratory Pro. We used standard instance (runtime), which have 12GB of RAM with NVIDIA NVIDIA Tesla P100 Graphic Cards having vRAM of 16GB.

5. Results and Discussion

6. Conclusions

7. Acknowledgements

Our project is part of Spring 2020 semester's Statistics 453⁷ (Introduction to Deep Learning and Generative Models) course of the University of Wisconsin-Madison. We specially appreciate with Prof. Raschka for all of his effort toward lectures and support toward our project.

8. Contributions

All of our team member decide to work on all steps and tasks together, as we believe it is important for everyone to at least understand how and what we are doing to achieve the common goal. However, as each individual has his/her own specialties and strength, some tasks has been designated to specific person.

As Hyecheol (Jerry) Jang has strength on understanding and fixing the codes written on python, he mainly worked on analyzing and fixing the baseline code. To do so, he went through the tensorflow API to understand what's going on inside the code and to remove depressed and unsupported code parts from the baseline code.

Alex studied similar studies and papers that are about creating music in various ways. He did research on the papers to find out which method worked and why specific methods worked better than the others. Also, he reviewed different models to find out which model would fit the best for the purpose of project.

Stella worked on finding the right MIDI files for the project. She went through various websites and databases to select the correct MIDI files that could be used for the purpose of our project. Also, she worked on understanding the code of the baseline model for the data processing part, to understand why the data was processed in such way, she had to study the MIDI file characteristics and music as well.

⁷Course Website: <http://pages.stat.wisc.edu/~sraschka/teaching/stat453-ss2020/>
GitHub Repository: <https://github.com/rasbt/stat453-deep-learning-ss20>

Even though we had some specific tasks assigned to each person, overall, as mentioned above, we helped each other and worked together on most of the tasks. Also, we divided the work on the presentation and the report evenly as well.

9. Codes

All the codes are posted on the GitHub repository.⁸ The repository ipynb file that used as an interactive workspace to run the code. All the data preprocessing codes are embedded on the notebook as we want to run those code only when needed. The model has been implemented in the python files, and it is called inside the notebook.

References

- [1] M. M. Association. General midi (gm 1). <https://www.midi.org/specifications-old/item/general-midi>.
- [2] G. Brunner, Y. Wang, R. Wattenhofer, and S. Zhao. Symbolic music genre transfer with cyclegan. In *2018 IEEE 30th International Conference on Tools with Artificial Intelligence (ICTAI)*, pages 786–793. IEEE, 2018.
- [3] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision*, pages 2223–2232, 2017.

⁸<https://github.com/hyecheol123/CycleGAN-Music-Style-Transfer>